

LoanTap Case Study

October 15, 2025

0.1 Introduction

LoanTap is at the forefront of offering tailored financial solutions to millennials. Their innovative approach seeks to harness data science for refining their credit underwriting process. The focus here is the Personal Loan segment. A deep dive into the dataset can reveal patterns in borrower behaviour and creditworthiness. Analyzing this dataset can provide crucial insights into the financial behaviours, spending habits and potential risk associated with each borrower. The insights gained can optimize loan disbursement, balancing customer outreach with risk management.

0.1.1 What is expected

Assuming you are a data scientist at LoanTap, you are tasked with analyzing the dataset to determine the creditworthiness of potential borrowers. Your ultimate objective is to build a logistic regression model, evaluate its performance, and provide actionable insights for the underwriting process.

0.1.2 Libraries

```
[1]: # libraries to analyze data
import numpy as np
import pandas as pd

# libraries to visualize data
import matplotlib.pyplot as plt
import seaborn as sns

import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

from sklearn.impute import KNNImputer
from sklearn.preprocessing import OneHotEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, \
    ConfusionMatrixDisplay, roc_curve, auc, precision_recall_curve, \
    average_precision_score
from sklearn.pipeline import make_pipeline

from imblearn.over_sampling import SMOTE
```

0.1.3 Data Loading

- Loading the data into Pandas dataframe for easily handling of data

```
[2]: # read the file into a pandas dataframe
customer_df = pd.read_csv("logistic_regression.csv")
df = customer_df
# look at the datatypes of the columns
print('*****')
print(df.info())
print('*****\n')
print('*****')
print(f'Shape of the dataset is {df.shape}')
print('*****\n')
print('*****')
print(f'Number of nan/null values in each column: \n{df.isna().sum()}')
print('*****\n')
print('*****')
print(f'Number of unique values in each column: \n{df.nunique()}')
print('*****\n')
print('*****')
print(f'Duplicate entries: \n{df.duplicated().value_counts()}')
print('*****')
```

```
*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt              396030 non-null  float64
1   term                   396030 non-null  object
2   int_rate                396030 non-null  float64
3   installment            396030 non-null  float64
4   grade                  396030 non-null  object
5   sub_grade              396030 non-null  object
6   emp_title              373103 non-null  object
7   emp_length             377729 non-null  object
8   home_ownership         396030 non-null  object
9   annual_inc             396030 non-null  float64
10  verification_status    396030 non-null  object
11  issue_d                396030 non-null  object
12  loan_status            396030 non-null  object
13  purpose                396030 non-null  object
14  title                  394274 non-null  object
15  dti                    396030 non-null  float64
16  earliest_cr_line       396030 non-null  object
17  open_acc               396030 non-null  float64
18  pub_rec                396030 non-null  float64
```

```

19  revol_bal          396030 non-null float64
20  revol_util         395754 non-null float64
21  total_acc          396030 non-null float64
22  initial_list_status 396030 non-null object
23  application_type    396030 non-null object
24  mort_acc           358235 non-null float64
25  pub_rec_bankruptcies 395495 non-null float64
26  address            396030 non-null object

```

dtypes: float64(12), object(15)

memory usage: 81.6+ MB

None

Shape of the dataset is (396030, 27)

Number of nan/null values in each column:

```

loan_amnt          0
term               0
int_rate           0
installment        0
grade              0
sub_grade           0
emp_title          22927
emp_length         18301
home_ownership      0
annual_inc          0
verification_status 0
issue_d             0
loan_status         0
purpose            0
title              1756
dti                 0
earliest_cr_line    0
open_acc            0
pub_rec             0
revol_bal           0
revol_util          276
total_acc           0
initial_list_status 0
application_type    0
mort_acc           37795
pub_rec_bankruptcies 535
address             0

```

dtype: int64

Number of unique values in each column:

loan_amnt	1397
term	2
int_rate	566
installment	55706
grade	7
sub_grade	35
emp_title	173105
emp_length	11
home_ownership	6
annual_inc	27197
verification_status	3
issue_d	115
loan_status	2
purpose	14
title	48816
dti	4262
earliest_cr_line	684
open_acc	61
pub_rec	20
revol_bal	55622
revol_util	1226
total_acc	118
initial_list_status	2
application_type	3
mort_acc	33
pub_rec_bankruptcies	9
address	393700

dtype: int64

Duplicate entries:

False 396030

Name: count, dtype: int64

```
[3]: # look at the top 5 rows
df.head()
```

```
[3]:   loan_amnt   term  int_rate  installment  grade  sub_grade  \
0   10000.0  36 months   11.44         329.48    B         B4
1    8000.0  36 months   11.99         265.68    B         B5
2   15600.0  36 months   10.49         506.97    B         B3
3    7200.0  36 months    6.49         220.65    A         A2
```

```
4      24375.0    60 months      17.27      609.33    C      C5
```

```

      emp_title emp_length home_ownership  annual_inc  ...  \
0      Marketing  10+ years      RENT    117000.0  ...
1    Credit analyst    4 years    MORTGAGE    65000.0  ...
2    Statistician    < 1 year      RENT    43057.0  ...
3    Client Advocate    6 years      RENT    54000.0  ...
4 Destiny Management Inc.    9 years    MORTGAGE    55000.0  ...

```

```

open_acc pub_rec revol_bal revol_util total_acc  initial_list_status  \
0     16.0     0.0   36369.0      41.8      25.0                      w
1     17.0     0.0   20131.0      53.3      27.0                      f
2     13.0     0.0   11987.0      92.2      26.0                      f
3      6.0     0.0    5472.0      21.5      13.0                      f
4     13.0     0.0   24584.0      69.8      43.0                      f

```

```

application_type  mort_acc  pub_rec_bankruptcies  \
0    INDIVIDUAL      0.0              0.0
1    INDIVIDUAL      3.0              0.0
2    INDIVIDUAL      0.0              0.0
3    INDIVIDUAL      0.0              0.0
4    INDIVIDUAL      1.0              0.0

```

```

                                address
0    0174 Michelle Gateway\r\nMendozaberg, OK 22690
1    1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113
2    87025 Mark Dale Apt. 269\r\nNew Sabrina, WV 05113
3      823 Reid Ford\r\nDelacruzside, MA 00813
4      679 Luna Roads\r\nGreggshire, VA 11650

```

[5 rows x 27 columns]

```
[4]: df[df.columns[10:20]]
```

```

[4]:      verification_status  issue_d  loan_status      purpose  \
0      Not Verified  Jan-2015    Fully Paid    vacation
1      Not Verified  Jan-2015    Fully Paid  debt_consolidation
2    Source Verified  Jan-2015    Fully Paid    credit_card
3      Not Verified  Nov-2014    Fully Paid    credit_card
4      Verified  Apr-2013  Charged Off    credit_card
...
396025    Source Verified  Oct-2015    Fully Paid  debt_consolidation
396026    Source Verified  Feb-2015    Fully Paid  debt_consolidation
396027      Verified  Oct-2013    Fully Paid  debt_consolidation
396028      Verified  Aug-2012    Fully Paid  debt_consolidation
396029      Verified  Jun-2010    Fully Paid  debt_consolidation

```

	title	dti	earliest_cr_line	open_acc	pub_rec	\
0	Vacation	26.24	Jun-1990	16.0	0.0	
1	Debt consolidation	22.05	Jul-2004	17.0	0.0	
2	Credit card refinancing	12.79	Aug-2007	13.0	0.0	
3	Credit card refinancing	2.60	Sep-2006	6.0	0.0	
4	Credit Card Refinance	33.95	Mar-1999	13.0	0.0	
...	...	...	...	...	...	
396025	Debt consolidation	15.63	Nov-2004	6.0	0.0	
396026	Debt consolidation	21.45	Feb-2006	6.0	0.0	
396027	pay off credit cards	17.56	Mar-1997	15.0	0.0	
396028	Loanforpayoff	15.88	Nov-1990	9.0	0.0	
396029	Toxic Debt Payoff	8.32	Sep-1998	3.0	0.0	

	revol_bal
0	36369.0
1	20131.0
2	11987.0
3	5472.0
4	24584.0
...	...
396025	1990.0
396026	43263.0
396027	32704.0
396028	15704.0
396029	4292.0

[396030 rows x 10 columns]

```
[5]: df[df.columns[20:]]
```

	revol_util	total_acc	initial_list_status	application_type	mort_acc	\
0	41.8	25.0	w	INDIVIDUAL	0.0	
1	53.3	27.0	f	INDIVIDUAL	3.0	
2	92.2	26.0	f	INDIVIDUAL	0.0	
3	21.5	13.0	f	INDIVIDUAL	0.0	
4	69.8	43.0	f	INDIVIDUAL	1.0	
...	...	...	...	...	...	
396025	34.3	23.0	w	INDIVIDUAL	0.0	
396026	95.7	8.0	f	INDIVIDUAL	1.0	
396027	66.9	23.0	f	INDIVIDUAL	0.0	
396028	53.8	20.0	f	INDIVIDUAL	5.0	
396029	91.3	19.0	f	INDIVIDUAL	NaN	

	pub_rec_bankruptcies	\
0	0.0	
1	0.0	
2	0.0	

```

3          0.0
4          0.0
...
396025     0.0
396026     0.0
396027     0.0
396028     0.0
396029     0.0

```

```

                                address
0      0174 Michelle Gateway\r\nMendozaberg, OK 22690
1      1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113
2      87025 Mark Dale Apt. 269\r\nNew Sabrina, WV 05113
3          823 Reid Ford\r\nDelacruzside, MA 00813
4          679 Luna Roads\r\nGreggshire, VA 11650
...
396025  12951 Williams Crossing\r\nJohnnyville, DC 30723
396026  0114 Fowler Field Suite 028\r\nRachelborough, ...
396027  953 Matthew Points Suite 414\r\nReedfort, NY 7...
396028  7843 Blake Freeway Apt. 229\r\nNew Michael, FL...
396029      787 Michelle Causeway\r\nBriannaton, AR 48052

```

[396030 rows x 7 columns]

```
[6]: df.describe()
```

```

[6]:      loan_amnt      int_rate      installment      annual_inc  \
count  396030.000000  396030.000000  396030.000000  3.960300e+05
mean    14113.888089    13.639400    431.849698    7.420318e+04
std      8357.441341     4.472157    250.727790    6.163762e+04
min       500.000000     5.320000     16.080000    0.000000e+00
25%      8000.000000    10.490000    250.330000    4.500000e+04
50%     12000.000000    13.330000    375.430000    6.400000e+04
75%     20000.000000    16.490000    567.300000    9.000000e+04
max     40000.000000    30.990000   1533.810000    8.706582e+06

      dti      open_acc      pub_rec      revol_bal  \
count  396030.000000  396030.000000  396030.000000  3.960300e+05
mean     17.379514    11.311153     0.178191    1.584454e+04
std     18.019092     5.137649     0.530671    2.059184e+04
min       0.000000     0.000000     0.000000    0.000000e+00
25%     11.280000     8.000000     0.000000    6.025000e+03
50%     16.910000    10.000000     0.000000    1.118100e+04
75%     22.980000    14.000000     0.000000    1.962000e+04
max     9999.000000    90.000000    86.000000    1.743266e+06

      revol_util      total_acc      mort_acc  pub_rec_bankruptcies

```

count	395754.000000	396030.000000	358235.000000	395495.000000
mean	53.791749	25.414744	1.813991	0.121648
std	24.452193	11.886991	2.147930	0.356174
min	0.000000	2.000000	0.000000	0.000000
25%	35.800000	17.000000	0.000000	0.000000
50%	54.800000	24.000000	1.000000	0.000000
75%	72.900000	32.000000	3.000000	0.000000
max	892.300000	151.000000	34.000000	8.000000

Insight

- There are 396030 entries with 27 columns
- There are 22927 null/missing values in emp_title, 18301 in emp_length, 1756 in title, 276 in revol_util, 37795 in mort_acc and 535 in pub_rec_bankruptcies
- There are no duplicates
- Based on the number of unique values in each column and their datatype, the columns term, grade, emp_length, home_ownership, verification_status, loan_status, purpose, initial_list_status and application_type can be converted to categorical datatype
- issue_d and earliest_cr_line has date values so can be converted to datetime
- The column sub_grade can be dropped as its info is already captured in column grade
- The columns term and emp_length have space, special characters, alphanumeric mix.
- Replace “36 months” with “short” and “60 months” with “long” in term column
- Remove “year/years”, replacing “< 1” with “0” and replacing “10+” with “10” in emp_length column. Then convert emp_length column to categorical datatype
- Extract zip code from address column and drop address column

```
[7]: import warnings

# Ignore FutureWarnings
warnings.simplefilter(action='ignore', category=FutureWarning)

# Convert to category
categorical_columns = ['term', 'grade', 'emp_length', 'home_ownership',
↳ 'verification_status', 'loan_status', 'purpose', 'initial_list_status',
↳ 'application_type']
df[categorical_columns] = df[categorical_columns].astype('category')

# Convert to datetime
df['issue_d'] = pd.to_datetime(df['issue_d'], format='%b-%Y')
df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'], format='%b-%Y')

# Drop "grade" column
df.drop(columns=['sub_grade'], inplace=True)
```



```

# Rename the values in 'term' column
df['term'].replace({' 36 months' : 'short', ' 60 months' : 'long'}, inplace=True)

# Rename employee_length column values
df['emp_length'] = df['emp_length'].replace({'< 1 year': '0 year'})
df['emp_length'] = df['emp_length'].str.replace(r'\D', '', regex=True)
df['emp_length'] = df['emp_length'].astype('category')

# Extract zip code from address
df['zip_code'] = df['address'].str[-5:].str.split().str[0].astype('category')
df.drop(columns='address', inplace=True)
print(f'Number of nan/null values in zip_code column: \n{df.zip_code.isna().sum()}')

df.info()

```

Number of nan/null values in zip_code column:

0

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 396030 entries, 0 to 396029

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
0	loan_amnt	396030 non-null	float64
1	term	396030 non-null	category
2	int_rate	396030 non-null	float64
3	installment	396030 non-null	float64
4	grade	396030 non-null	category
5	emp_title	373103 non-null	object
6	emp_length	377729 non-null	category
7	home_ownership	396030 non-null	category
8	annual_inc	396030 non-null	float64
9	verification_status	396030 non-null	category
10	issue_d	396030 non-null	datetime64[ns]
11	loan_status	396030 non-null	category
12	purpose	396030 non-null	category
13	title	394274 non-null	object
14	dti	396030 non-null	float64
15	earliest_cr_line	396030 non-null	datetime64[ns]
16	open_acc	396030 non-null	float64
17	pub_rec	396030 non-null	float64
18	revol_bal	396030 non-null	float64
19	revol_util	395754 non-null	float64
20	total_acc	396030 non-null	float64
21	initial_list_status	396030 non-null	category

```

22 application_type      396030 non-null  category
23 mort_acc              358235 non-null  float64
24 pub_rec_bankruptcies  395495 non-null  float64
25 zip_code              396030 non-null  category
dtypes: category(10), datetime64[ns](2), float64(12), object(2)
memory usage: 52.1+ MB

```

0.1.4 Handling null values

- emp_title has 22927, emp_length has 18301, title has 1756, revol_util has 276, mort_acc has 37795 and pub_rec_bankruptcies has 535 null values
- Columns emp_title and title can be dropped as they would not have an effect on the loan approval
- Null values in revol_util and pub_rec_bankruptcies are small in number and hence can be dropped

```

[8]: df.drop(columns=['emp_title', 'title'], inplace=True)
df.dropna(subset = ['revol_util', 'pub_rec_bankruptcies'], inplace=True)


```

- Let us check the distribution of remaining features before deciding on how to handle the null values

```

[9]: fig, axs = plt.subplots(1,2, figsize=(10,4))
sns.histplot(ax = axs[0], data=df, x = 'emp_length')
sns.histplot(ax = axs[1], data=df, x = 'mort_acc')
fig.tight_layout()
plt.show()

```



- mort_acc is the number of mortgage accounts out of the total_acc. So I will replace the null values with “mode” of mort_acc for different total_acc
- For emp_length, I will use knn imputer to fill the missing data

```
[10]: mode_mort_acc_df = df.groupby('total_acc')['mort_acc'].agg(lambda x: pd.Series.
    ↪mode(x)[0])
```

```
def fill_mort_acc(total_acc, mort_acc):
    if np.isnan(mort_acc):
        return mode_mort_acc_df[total_acc]
    else:
        return mort_acc

df['mort_acc'] = df.apply(lambda x:
    ↪fill_mort_acc(x['total_acc'],x['mort_acc']), axis=1)
```

```
[11]: imputer = KNNImputer(n_neighbors=5)
df['emp_length']=imputer.fit_transform(df[['emp_length']])
df['emp_length'] = df['emp_length'].astype('category')
```

```
[15]: df.isna().sum()
```

```
[15]: loan_amnt          0
term                    0
int_rate                0
installment            0
grade                  0
emp_length              0
home_ownership          0
annual_inc              0
verification_status     0
issue_d                 0
loan_status             0
purpose                 0
dti                     0
earliest_cr_line        0
open_acc                0
pub_rec                 0
revol_bal               0
revol_util              0
total_acc               0
initial_list_status      0
application_type         0
mort_acc                0
pub_rec_bankruptcies     0
zip_code                 0
dtype: int64
```

- There are no null values now

0.1.5 Detecting outliers

Outliers for every continuous variable

```
[16]: # helper function to detect outliers using IQR method
def detectOutliers_iqr(df):
    q1 = df.quantile(0.25)
    q3 = df.quantile(0.75)
    iqr = q3-q1
    lower_outliers = df[df<(q1-1.5*iqr)]
    higher_outliers = df[df>(q3+1.5*iqr)]
    return lower_outliers, higher_outliers

[17]: # helper function to detect outliers using standard deviation method
def detectOutliers_std(df):
    mean = df.mean()
    std = df.std()
    upper_limit = mean+(3*std)
    lower_limit = mean-(3*std)
    lower_outliers = df[df<lower_limit]
    higher_outliers = df[df>upper_limit]
    return lower_outliers, higher_outliers

[18]: numerical_columns = df.select_dtypes(include=np.number).columns
column_outlier_dictionary = {}
for column in numerical_columns:
    lower_outliers, higher_outliers = detectOutliers_iqr(df[column])
    column_outlier_dictionary[column] = [lower_outliers, higher_outliers]
    #print('*'*50)
    #print(f'Outliers of \'{column}\'' column are:')
    #print("Lower outliers:\n", lower_outliers)
    #print("Higher outliers:\n", higher_outliers)
    #print('*'*50, end="\n")

[19]: for key, value in column_outlier_dictionary.items():
    print(f'The column \'{key}\'' has {len(value[0]) + len(value[1])} outliers')
```

```
The column 'loan_amnt' has 190 outliers
The column 'int_rate' has 3144 outliers
The column 'installment' has 11114 outliers
The column 'annual_inc' has 16649 outliers
The column 'dti' has 275 outliers
The column 'open_acc' has 10297 outliers
The column 'pub_rec' has 57730 outliers
The column 'revol_bal' has 21205 outliers
The column 'revol_util' has 12 outliers
The column 'total_acc' has 8491 outliers
The column 'mort_acc' has 6837 outliers
The column 'pub_rec_bankruptcies' has 45111 outliers
```

```
[81]: df.columns = df.columns.str.strip()
numerical_columns = [col.strip() for col in numerical_columns if col in df.
↳ columns]


# Setup layout
num_cols = 6
num_rows = int(np.ceil(len(numerical_columns)/num_cols))
fig, axs = plt.subplots(num_rows, num_cols, figsize=(15, num_rows*2.5))

# Create a blue gradient palette
blues = sns.color_palette("Blues", n_colors=len(numerical_columns))

# Plot each numerical column
for idx, col in enumerate(numerical_columns):
    ax = plt.subplot(num_rows, num_cols, idx+1)
    sns.boxplot(ax=ax, data=df, y=col, color=blues[idx])
    ax.set_title(col, fontsize=10)

# Remove any empty subplots
for i in range(len(numerical_columns), num_rows*num_cols):
    fig.delaxes(axs.flatten()[i])

plt.tight_layout()
plt.show()
```



```
[22]: for key, value in column_outlier_dictionary.items():
print(f'The column \'{key}\'' has {len(value[0]) + len(value[1])} outliers')
```

The column 'loan_amnt' has 184 outliers
The column 'int_rate' has 754 outliers
The column 'installment' has 5042 outliers
The column 'annual_inc' has 3190 outliers
The column 'dti' has 12 outliers
The column 'open_acc' has 4873 outliers

The column 'revol_bal' has 4771 outliers
The column 'revol_util' has 16 outliers
The column 'total_acc' has 3396 outliers
The column 'mort_acc' has 6837 outliers

0.1.6 Remove the outliers

- Based on the boxplot, the number of outliers using IQR method and standard deviation method, I will remove the outliers using the standard deviation method except for columns pub_rec and pub_rec_bankruptcies which will be removed based on manual check.

```
[23]: remove_outliers = True
      if True == remove_outliers:
          master_index = pd.core.indexes.base.Index([])
          for key, value in column_outlier_dictionary.items():
              lower_outliers = value[0]
              higher_outliers = value[1]
              master_index = master_index.union(lower_outliers.index).
              ↪union(higher_outliers.index)
          df.drop(master_index, inplace=True)
      else:
          print('Not removing any outliers')
```

```
[24]: df['pub_rec'].value_counts()
```


```
[24]: pub_rec
0.0    315552
1.0     47129
2.0      5107
3.0     1424
4.0      481
5.0      218
6.0      108
7.0       47
8.0       31
10.0      11
9.0       10
11.0       6
13.0       4
12.0       4
19.0       2
40.0       1
17.0       1
86.0       1
24.0       1
15.0       1
Name: count, dtype: int64
```

```
[25]: df['pub_rec_bankruptcies'].value_counts()
```

```
[25]: pub_rec_bankruptcies
0.0    327200
1.0    40774
2.0     1716
3.0     332
4.0      75
5.0     30
6.0      6
7.0      4
8.0      2
Name: count, dtype: int64
```

```
[26]: df = df[df['pub_rec'] < 11]
df = df[df['pub_rec_bankruptcies'] < 6]
```

```
[27]: numerical_columns = df.select_dtypes(include=np.number).columns
num_cols = 6
num_rows = int(np.ceil(len(numerical_columns)/num_cols))
fig, axs = plt.subplots(num_rows, num_cols, figsize=(10,15))
for idx in range(len(numerical_columns)):
    ax = plt.subplot(num_rows, num_cols, idx+1)
    sns.boxplot(ax = ax, data=df, y = numerical_columns[idx])
plt.tight_layout()
plt.show()
```




```
[28]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 370106 entries, 0 to 396029
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             370106 non-null float64
1   term                  370106 non-null category
2   int_rate              370106 non-null float64
3   installment           370106 non-null float64
4   grade                 370106 non-null category
5   emp_length            370106 non-null category
6   home_ownership        370106 non-null category
7   annual_inc            370106 non-null float64
8   verification_status   370106 non-null category
9   issue_d               370106 non-null datetime64[ns]
10  loan_status           370106 non-null category
11  purpose               370106 non-null category
12  dti                   370106 non-null float64
13  earliest_cr_line      370106 non-null datetime64[ns]
14  open_acc              370106 non-null float64
15  pub_rec               370106 non-null float64
16  revol_bal             370106 non-null float64
17  revol_util            370106 non-null float64
18  total_acc             370106 non-null float64
19  initial_list_status    370106 non-null category
20  application_type       370106 non-null category
21  mort_acc              370106 non-null float64
22  pub_rec_bankruptcies   370106 non-null float64
23  zip_code              370106 non-null category
dtypes: category(10), datetime64[ns](2), float64(12)
memory usage: 45.9 MB
```

Insights


- The number of columns reduced to 370106 from the original 396030 rows

0.1.7 Univariate analysis

Numerical Variables

```
[30]: num_cols = 2
num_rows = int(np.ceil(len(numerical_columns)/num_cols))
fig, axs = plt.subplots(num_rows,num_cols,figsize=(10,15))
for idx in range(len(numerical_columns)):
    ax = plt.subplot(num_rows, num_cols, idx+1)
```

```
sns.histplot(ax = ax, data=df, x = numerical_columns[idx], kde=True)
plt.tight_layout()
plt.show()
```




Insight

- Most of the features are right skewed except for revol_util which is slightly left skewed

0.1.8 Categorical Variables

```
[31]: categorical_columns = df.select_dtypes(include='category').columns
num_cols = 2
num_rows = int(np.ceil(len(categorical_columns)/num_cols))
fig, axs = plt.subplots(num_rows, num_cols, figsize=(12,30))
for idx in range(len(categorical_columns)):
    ax = plt.subplot(num_rows, num_cols, idx+1)
    sns.countplot(ax = ax, data=df, y = categorical_columns[idx],
        ↪stat='percent')
plt.tight_layout()
plt.show()
```




Insights

- Most of the loan is taken for a short term which is for 3 years
- Maximum loan are assigned the grade B followed by grade C
- Maximum loan are taken by borrowers whose employment tenure is more than 10 years
- Maximum loan are taken by borrowers whose house is mortgaged followed by those whose who are in rented house
- Most of the loan is fully_paid
- Most of the loan is taken for debt consolidation followed by credit card
- Majority of the loan is taken by individuals

0.1.9 Bivariate analysis


```
[32]: sns.boxplot(data=df, x='loan_status', y='loan_amnt')  
plt.show()
```



Insights

- The median of the loan amount slightly higher for loans which were charged off

```
[33]: fig, axs = plt.subplots(2,2,figsize=(10,6))
sns.scatterplot(ax=axs[0,0], data=df, x='loan_amnt', y='installment')
sns.scatterplot(ax=axs[0,1], data=df, x='loan_amnt', y='annual_inc')
sns.scatterplot(ax=axs[1,0], data=df, x='annual_inc', y='dti')
sns.scatterplot(ax=axs[1,1], data=df, x='total_acc', y='open_acc')
fig.tight_layout()
plt.show()
```




Insights

- It is very clear from the plot that, in general, the installment increases as the loan amount increases
- It is very obvious that people with higher income can afford to take higher loan
- As income increases debt to income ratio reduces
- Most of the borrower's accounts are active accounts

```
[34]: fig, axs = plt.subplots(2,2,figsize=(10,6))
sns.histplot(ax=axs[0,0], data=df, x='pub_rec', hue='loan_status')
sns.histplot(ax=axs[0,1], data=df, x='mort_acc', hue='loan_status')
sns.histplot(ax=axs[1,0], data=df, x='pub_rec_bankruptcies', hue='loan_status')
fig.tight_layout()
```

```
plt.show()
```



Insights

- I will group pub_rec into 2 groups: 0 for having 0 negative records and 1 for having more than 0 negative records
- Same for mort_acc and pub_rec_bankruptcies

```
[35]: group_0_list = [0.0]
pub_rec_list = list(df['pub_rec'].explode().unique())
group_1_list = list(set(pub_rec_list) - set(group_0_list))
df['any_neg_rec'] = df['pub_rec'].replace(group_0_list, 0)
df['any_neg_rec'] = df['any_neg_rec'].replace(group_1_list, 1)
df['any_neg_rec'] = df['any_neg_rec'].astype('category')

group_0_list = [0.0]
pub_rec_list = list(df['mort_acc'].explode().unique())
group_1_list = list(set(pub_rec_list) - set(group_0_list))
df['any_mort'] = df['mort_acc'].replace(group_0_list, 0)
df['any_mort'] = df['any_mort'].replace(group_1_list, 1)
df['any_mort'] = df['any_mort'].astype('category')


group_0_list = [0.0]
pub_rec_list = list(df['pub_rec_bankruptcies'].explode().unique())
group_1_list = list(set(pub_rec_list) - set(group_0_list))
```



```
df['any_bankruptcies'] = df['pub_rec_bankruptcies'].replace(group_0_list, 0)
df['any_bankruptcies'] = df['any_bankruptcies'].replace(group_1_list, 1)
df['any_bankruptcies'] = df['any_bankruptcies'].astype('category')

df.drop(columns = ['pub_rec', 'mort_acc', 'pub_rec_bankruptcies'], inplace=True)
```

```
[36]: fig, axs = plt.subplots(2,2,figsize=(10,6))
sns.histplot(ax=axs[0,0], data=df, x='any_neg_rec', hue='loan_status',
             stat="proportion", multiple="fill")
sns.histplot(ax=axs[0,1], data=df, x='any_mort', hue='loan_status',
             stat="proportion", multiple="fill")
sns.histplot(ax=axs[1,0], data=df, x='any_bankruptcies', hue='loan_status',
             stat="proportion", multiple="fill")
fig.tight_layout()
plt.show()
```




```
[37]: fig, axs = plt.subplots(3,2,figsize=(12,10))
sns.histplot(ax=axs[0,0], data=df, x='term', hue='loan_status',
             stat="proportion", multiple="fill")
sns.histplot(ax=axs[0,1], data=df, x='home_ownership', hue='loan_status',
             stat="proportion", multiple="fill")
sns.histplot(ax=axs[1,0], data=df, x='verification_status', hue='loan_status',
             stat="proportion", multiple="fill")
sns.histplot(ax=axs[1,1], data=df, x='initial_list_status', hue='loan_status',
             stat="proportion", multiple="fill")
```


```

sns.histplot(ax=axis[2,0], data=df, x='application_type', hue='loan_status',
             stat="proportion", multiple="fill")
fig.tight_layout()
plt.show()

fig, axis = plt.subplots(4,1,figsize=(15,15))
sns.histplot(ax=axis[0], data=df, x='grade', hue='loan_status',
             stat="proportion", multiple="fill")
sns.histplot(ax=axis[1], data=df, x='emp_length', hue='loan_status',
             stat="proportion", multiple="fill")
sns.histplot(ax=axis[2], data=df, y='purpose', hue='loan_status',
             stat="proportion", multiple="fill")
sns.histplot(ax=axis[3], data=df, x='zip_code', hue='loan_status',
             stat="proportion", multiple="fill")
fig.tight_layout()
plt.show()

```





Insights


- Having a negative or bankruptcy record doesn't seem to impact the loan getting paid back or not
- Loan taken for short term, i.e. 3 years are most likely to be fully paid back
- Loan taken by people whose house ownership is of type any is on are paid back 100%.
- Surprisingly, loans which are not verified are more likely to be paid back
- Initial list status also does not seem to impact loan status
- Loan taken as joint application type are more likely to be paid back
- Loans with grade A and B are more likely to get paid back

- Employment duration does not seem to impact loan status
- Loan taken for wedding are more likely to be paid back
- Loan taken by people with zip code 00813 and 05113 are fully paid back whereas loan taken by people with zip code 11650, 86630 and 93700 are all charged off

```
[39]: df.drop(columns=['any_neg_rec', 'any_bankruptcies', 'initial_list_status', 'emp_length'], inplace=True)
```

0.1.10 Multivariate analysis

```
[44]: fig, ax = plt.subplots(figsize=(8,6))
sns.heatmap(df.select_dtypes(include=np.number).corr(), annot=True, linewidth=0.5, cmap = "Blues", ax=ax)
ax.xaxis.tick_top()
plt.xticks(rotation=30, ha='left')
plt.show()
```



Insights

- loan amount is highly correlated with installment
- There is good correlation between loan amount - annual income, loan amount - revol balance, installment - annual income, installment - revol balance, open account - total account

0.1.11 Data Preprocessing

[45]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 370106 entries, 0 to 396029
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             370106 non-null float64
1   term                  370106 non-null category
2   int_rate              370106 non-null float64
3   installment           370106 non-null float64
4   grade                 370106 non-null category
5   home_ownership        370106 non-null category
6   annual_inc            370106 non-null float64
7   verification_status   370106 non-null category
8   issue_d               370106 non-null datetime64[ns]
9   loan_status           370106 non-null category
10  purpose               370106 non-null category
11  dti                   370106 non-null float64
12  earliest_cr_line      370106 non-null datetime64[ns]
13  open_acc              370106 non-null float64
14  revol_bal             370106 non-null float64
15  revol_util            370106 non-null float64
16  total_acc             370106 non-null float64
17  application_type      370106 non-null category
18  zip_code              370106 non-null category
19  any_mort              370106 non-null category
dtypes: category(9), datetime64[ns](2), float64(9)
memory usage: 37.1 MB
```

- The date features will not have an impact on the loan status, so i will drop `issue_d` and `earliest_cr_line` columns

[46]: `df.drop(columns=['issue_d', 'earliest_cr_line'], inplace=True)`

0.1.12 Multicollinearity Check

```
[47]: features_df = df.drop(columns=['loan_status']) # Drop target column
features_df.select_dtypes(include='category').columns
features_df = features_df.drop(columns=features_df.
    ↪select_dtypes(include='category').columns) # Drop category columns
```

```

features_df = sm.add_constant(features_df) # Adding a constant column for the
↳intercept
vif_df = pd.DataFrame()
vif_df['Features'] = features_df.columns
vif_df['VIF'] = [variance_inflation_factor(features_df.values, idx) for idx in
↳range(len(features_df.columns))]
vif_df['VIF'] = round(vif_df['VIF'], 2)
vif_df = vif_df.sort_values(by='VIF', ascending=False)
vif_df

```

```

[47]:
      Features    VIF
0      const  24.91
1  loan_amnt  11.58
3  installment  11.04
6   open_acc   2.00
9  total_acc   1.87
7  revol_bal   1.76
4  annual_inc   1.62
8  revol_util   1.47
5         dti   1.40
2   int_rate   1.23

```

Insights

- loan amount is highly correlated with installment which is also shown here by high VIF values.
I will drop installment

```

[48]: features_df = features_df.drop(columns=['installment'])
features_df = sm.add_constant(features_df) # Adding a constant column for the
↳intercept
vif_df = pd.DataFrame()
vif_df['Features'] = features_df.columns
vif_df['VIF'] = [variance_inflation_factor(features_df.values, idx) for idx in
↳range(len(features_df.columns))]
vif_df['VIF'] = round(vif_df['VIF'], 2)
vif_df = vif_df.sort_values(by='VIF', ascending=False)
vif_df

```

```

[48]:
      Features    VIF
0      const  24.76
5   open_acc   2.00
8  total_acc   1.86
6  revol_bal   1.75
3  annual_inc   1.62
1  loan_amnt   1.49
7  revol_util   1.46
4         dti   1.40

```

```
2    int_rate    1.22
```

Insights

- Based on the above VIF scores, I can conclude that there are no more multicollinear numerical features
- I will drop installment from the dataframe

```
[49]: df.drop(columns=['installment'], inplace=True)
```

0.1.13 Encode categorical variables

```
[50]: final_df = df.copy()
final_df.reset_index(inplace=True, drop=True)
```

Seperate out target and feature columns

```
[52]: X = final_df.drop(columns=['loan_status'])
y = final_df['loan_status']
```

Encode target variable

```
[53]: y = y.replace({'Fully Paid': 0, 'Charged Off': 1}).astype(int)
```

Encode features with just 2 classes as 0 or 1

```
[54]: X['term'] = X['term'].replace({'short': 0, 'long': 1}).astype('int8')
X['any_mort'] = X['any_mort'].astype('int8')
```

One-Hot-Encoding for remaining categorical features

```
[55]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 370106 entries, 0 to 370105
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             370106 non-null  float64
1   term                  370106 non-null  int8
2   int_rate              370106 non-null  float64
3   grade                 370106 non-null  category
4   home_ownership        370106 non-null  category
5   annual_inc            370106 non-null  float64
6   verification_status   370106 non-null  category
7   purpose               370106 non-null  category
8   dti                   370106 non-null  float64
9   open_acc              370106 non-null  float64
10  revol_bal             370106 non-null  float64
11  revol_util            370106 non-null  float64
```

```

12 total_acc          370106 non-null float64
13 application_type   370106 non-null category
14 zip_code           370106 non-null category
15 any_mort           370106 non-null int8
dtypes: category(6), float64(8), int8(2)
memory usage: 25.4 MB

```

```
[56]: categorical_columns = X.select_dtypes(include='category').columns
categorical_columns
```

```
[56]: Index(['grade', 'home_ownership', 'verification_status', 'purpose',
          'application_type', 'zip_code'],
          dtype='object')
```

```
[57]: encoder = OneHotEncoder(sparse_output=False)
encoded_data = encoder.fit_transform(X[categorical_columns])
encoded_df = pd.DataFrame(encoded_data, columns = encoder.
    ↳get_feature_names_out(categorical_columns))
X = pd.concat([X, encoded_df], axis=1)
X.drop(columns = categorical_columns, inplace=True)
X.head()
```

```
[57]:
```

	loan_amnt	term	int_rate	annual_inc	dti	open_acc	revol_bal	\
0	10000.0	0	11.44	117000.0	26.24	16.0	36369.0	
1	8000.0	0	11.99	65000.0	22.05	17.0	20131.0	
2	15600.0	0	10.49	43057.0	12.79	13.0	11987.0	
3	7200.0	0	6.49	54000.0	2.60	6.0	5472.0	
4	24375.0	1	17.27	55000.0	33.95	13.0	24584.0	

	revol_util	total_acc	any_mort	...	zip_code_00813	zip_code_05113	\
0	41.8	25.0	0	...	0.0	0.0	
1	53.3	27.0	1	...	0.0	1.0	
2	92.2	26.0	0	...	0.0	1.0	
3	21.5	13.0	0	...	1.0	0.0	
4	69.8	43.0	1	...	0.0	0.0	

	zip_code_11650	zip_code_22690	zip_code_29597	zip_code_30723	\
0	0.0	1.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	1.0	0.0	0.0	0.0	

	zip_code_48052	zip_code_70466	zip_code_86630	zip_code_93700
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0

4	0.0	0.0	0.0	0.0
---	-----	-----	-----	-----

[5 rows x 53 columns]

0.1.14 Train-test split

```
[58]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=0)
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[58]: ((296084, 53), (74022, 53), (296084,), (74022,))
```

0.1.15 Perform data normalization/standardization

Data normalization/standardization is required so that features with higher scales do not dominate the model's performance. Hence all features should have same scale

Data before scaling

```
[59]: X_train.head()
```

```
[59]:
```

	loan_amnt	term	int_rate	annual_inc	dti	open_acc	revol_bal	\
133405	27000.0	1	16.29	82302.0	25.52	13.0	12014.0	
365868	6000.0	0	18.55	45000.0	19.37	8.0	3219.0	
71124	8975.0	0	9.71	65000.0	7.98	10.0	3932.0	
33923	9600.0	0	6.62	58000.0	25.01	10.0	57236.0	
30512	18000.0	0	11.53	75000.0	8.50	9.0	9916.0	

	revol_util	total_acc	any_mort	...	zip_code_00813	zip_code_05113	\
133405	48.6	29.0	1	...	0.0	0.0	
365868	73.2	11.0	0	...	0.0	0.0	
71124	34.5	58.0	0	...	0.0	0.0	
33923	36.3	19.0	1	...	0.0	0.0	
30512	35.8	17.0	0	...	0.0	0.0	

	zip_code_11650	zip_code_22690	zip_code_29597	zip_code_30723	\
133405	0.0	0.0	0.0	0.0	
365868	0.0	0.0	0.0	0.0	
71124	0.0	0.0	0.0	0.0	
33923	0.0	0.0	0.0	0.0	
30512	0.0	0.0	1.0	0.0	

	zip_code_48052	zip_code_70466	zip_code_86630	zip_code_93700
133405	0.0	1.0	0.0	0.0
365868	0.0	0.0	0.0	1.0
71124	1.0	0.0	0.0	0.0
33923	1.0	0.0	0.0	0.0
30512	0.0	0.0	0.0	0.0

[5 rows x 53 columns]

```
[60]: min_max_scaler = MinMaxScaler()
# Fit min_max_scaler to training data
min_max_scaler.fit(X_train)
# Scale the training and testing data
X_train = pd.DataFrame(min_max_scaler.transform(X_train), columns=X_train.
↪columns)
X_test = pd.DataFrame(min_max_scaler.transform(X_test), columns=X_test.columns)
```

Data after scaling

```
[61]: X_train.head()
```

```
[61]:   loan_amnt  term  int_rate  annual_inc    dti  open_acc  revol_bal  \
0   0.697828   1.0   0.506230   0.308276  0.372229    0.48   0.154818
1   0.144832   0.0   0.610521   0.161417  0.282526    0.28   0.041481
2   0.223173   0.0   0.202584   0.240157  0.116394    0.36   0.050669
3   0.239631   0.0   0.059991   0.212598  0.364790    0.36   0.737568
4   0.460829   0.0   0.286571   0.279528  0.123979    0.32   0.127782

   revol_util  total_acc  any_mort  ...  zip_code_00813  zip_code_05113  \
0   0.394481   0.457627    1.0  ...             0.0             0.0
1   0.594156   0.152542    0.0  ...             0.0             0.0
2   0.280032   0.949153    0.0  ...             0.0             0.0
3   0.294643   0.288136    1.0  ...             0.0             0.0
4   0.290584   0.254237    0.0  ...             0.0             0.0

   zip_code_11650  zip_code_22690  zip_code_29597  zip_code_30723  \
0             0.0             0.0             0.0             0.0
1             0.0             0.0             0.0             0.0
2             0.0             0.0             0.0             0.0
3             0.0             0.0             0.0             0.0
4             0.0             0.0             1.0             0.0

   zip_code_48052  zip_code_70466  zip_code_86630  zip_code_93700
0             0.0             1.0             0.0             0.0
1             0.0             0.0             0.0             1.0
2             1.0             0.0             0.0             0.0
3             1.0             0.0             0.0             0.0
4             0.0             0.0             0.0             0.0
```

[5 rows x 53 columns]

Check for imbalance in target class

```
[62]: y_train.value_counts(normalize=True)*100
```

```
[62]: loan_status
0      80.249186
1      19.750814
Name: proportion, dtype: float64
```

- We can see a clear imbalance in the target class with 1 being ~20% and 0 being ~80%. Hence, I will use SMOTE to fix this imbalance

```
[63]: sm = SMOTE(random_state=0)
X_train, y_train = sm.fit_resample(X_train, y_train)
y_train.value_counts(normalize=True)*100
```

```
[63]: loan_status
0      50.0
1      50.0
Name: proportion, dtype: float64
```

0.1.16 Build Logistic Regression model

Train the model


```
[64]: model = LogisticRegression(solver='lbfgs', max_iter=300)
model.fit(X_train, y_train)
y_train_pred = model.predict(X_train)
```

- Classification metrics and confusion matrix for Training data

```
[65]: print(classification_report(y_train, y_train_pred))
```

	precision	recall	f1-score	support
0	0.81	0.80	0.80	237605
1	0.80	0.82	0.81	237605
accuracy			0.81	475210
macro avg	0.81	0.81	0.81	475210
weighted avg	0.81	0.81	0.81	475210

```
[66]: # Confusion Matrix
cm = confusion_matrix(y_train, y_train_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```



Model prediction

```
[67]: y_test_pred = model.predict(X_test)
```


Classification metrics and confusion matrix for Testing data

```
[68]: print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	0.95	0.80	0.86	59619
1	0.49	0.81	0.61	14403
accuracy			0.80	74022
macro avg	0.72	0.80	0.74	74022
weighted avg	0.86	0.80	0.82	74022

```
[69]: # Confusion Matrix
cm = confusion_matrix(y_test, y_test_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
```

```
plt.title('Confusion Matrix')
plt.show()
```




Insights

- Recall is high indicating that the model is able to identify 80% of the actual defaulters and 80% of non-defaulting customers
- Precision for class 1 (defaulters) is low. Of all the predicted defaulters, only 50% are actual defaulters
- With this model there is a risk of denying loans to deserving customers due to low precision score for defaulters

```
[70]: feature_imp = pd.DataFrame({'Columns':X_train.columns, 'Coefficients':model.
    ↪coef_[0]}).round(2).sort_values('Coefficients', ascending=False)

plt.figure(figsize=(8,8))
sns.barplot(data=feature_imp, y = 'Columns', x = 'Coefficients')
plt.title("Feature Importance for Model")
plt.yticks(fontsize=8)
```

```
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



Insights

- The features zip_code_29597, zip_code_05113, zip_code_00813, annual_inc and application_type_joint have got high positive weightage and features zip_code_86630, zip_code_11650, zip_code_93700, dti and open_acc have got high negative weightage indicating their major contribution towards target variable


0.1.17 ROC and AUC

```
[72]: # Predict probabilities for the test set
probs = model.predict_proba(X_test)[:,-1]

# Compute the false positive rate, true positive rate, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, probs)

# Compute the area under the ROC curve
roc_auc = auc(fpr, tpr)

# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
        roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



Insights


- ROC curve illustrates the trade off between TPR(True Positive Rate) and FPR(False Positive Rate) for various thresholds
- The AU-ROC value of 0.91 signifies that the model is able to differentiate well between the two classes
- Let us also look at PR Curve(Precision Recall Curve)

0.1.18 Precision Recall Curve

```
[73]: precision, recall, thr = precision_recall_curve(y_test, probs)

# Area under Precision Recall Curve
apc = average_precision_score(y_test, probs)

# Plot the precision-recall curve
plt.plot(recall, precision, marker='.', label='PR curve (area = %0.2f)' % apc)
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc="lower left")
plt.show()
```


Insights

- PR curve illustrates the trade off between Precision and Recall for various thresholds
- The model has a AU-PRC value of 0.78 which is not that high. It is better than the random model which has a AU-PRC value of 0.5.
- This clearly indicates that we simply cannot conclude on the model's performance from just the ROC curve.

Find the threshold where precision and recall meet

```
[74]: plt.figure()
plt.plot(thr,precision[0:len(thr)],label='precision',color='blue')
plt.plot(thr,recall[0:len(thr)],label='recall',color='orange')
intersection_thr = thr[np.where(precision == recall)[0][0]].round(4)
plt.axvline(intersection_thr, linestyle='--', color='red')
plt.text(intersection_thr, 0.01, str(intersection_thr), ha='left', color='red')
plt.title("Precision-recall curve at different thresholds")
plt.xlabel("Threshold values")
plt.ylabel("Precision and Recall values")
plt.legend(loc="upper right")
```

```
plt.grid()
plt.show()
```




```
[75]: y_pred = model.predict_proba(X_test)[: ,1]
threshold_considered = intersection_thr
y_pred_custom = (y_pred>threshold_considered).astype('int')
print(classification_report(y_test,y_pred_custom))
```

	precision	recall	f1-score	support
0	0.92	0.92	0.92	59619
1	0.65	0.65	0.65	14403
accuracy			0.87	74022
macro avg	0.79	0.79	0.79	74022
weighted avg	0.87	0.87	0.87	74022

```
[76]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred_custom)
```

```
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```



Insights

- With the new threshold, precision for class 1 (defaulters) has increased at a cost of decrease in Recall
- The overall F1-score and accuracy has increased

0.2 Actionable Insights:

1. Loan Payment Status Distribution:

- 80% of the customers have fully paid their loan, while 20% are defaulters.
- The imbalance between the classes (fully paid vs. charged off) can impact model performance, especially for the minority class (charged off).

2. Model Performance:

- F1-Score for Negative Class (Fully Paid): The model achieves a 94% F1-score for the negative class, indicating that it performs well at predicting customers who will fully pay the loan.
 - F1-Score for Positive Class (Charged Off): The model achieves 62% F1-score for the positive class, indicating lower performance in predicting defaulters.
 - This shows that the model is better at identifying customers who will pay fully, but there's room for improvement in predicting defaulters.
3. Cross-Validation and Testing Accuracy:
 - The cross-validation accuracy and test accuracy are almost identical, which suggests that the model is generalizing well to unseen data and is not overfitting.
 4. ROC-AUC Score:
 - The ROC AUC score is 0.73, indicating that the model is correctly classifying around 73% of the instances.
 - While this is a good performance, there is still room for improvement in the model's ability to distinguish between the classes.
 5. Precision-Recall Curve:
 - The precision-recall curve reveals the trade-off between precision and recall as the threshold changes.
 - Adjusting the threshold can help achieve the optimal balance depending on the business need—higher precision leads to fewer false positives (charged off classified as fully paid), while higher recall increases the likelihood of identifying defaulters.
 6. Impact of Data Balancing:
 - After balancing the dataset, there was a significant change in both precision and recall scores for both classes.
 - Balancing the classes has helped improve the model's ability to predict both outcomes more fairly.
 7. Logistic Regression Accuracy:
 - The accuracy of the Logistic Regression classifier on the test set is 0.891, which indicates a solid performance that's not due to chance.

0.3 Recommendations:

1. Improving Model Performance for Defaulters:
 - Since the model performs better for the negative class (fully paid), there should be an emphasis on improving predictions for the positive class (defaulters).
 - Consider using techniques like oversampling (SMOTE) or undersampling to address the class imbalance and improve model sensitivity towards the minority class.
2. Advanced Models:

- To further enhance performance, especially for predicting defaulters, consider using more complex models like Random Forests, Gradient Boosting, or XGBoost.
 - These models are typically better at handling class imbalances and complex decision boundaries.
3. Hyperparameter Tuning:
- Conduct hyperparameter tuning using techniques such as GridSearchCV or RandomizedSearchCV to optimize model parameters for better performance, particularly on the defaulter class.
 - Adjust parameters like regularization strength in Logistic Regression or tree depths in decision tree-based models to improve the model's predictive power.
4. Feature Engineering:
- Investigate additional feature engineering to create more informative features that could help improve the model's predictive accuracy, such as aggregating payment behavior over time or including customer demographic information.
5. Collect More Data:
- Data quantity plays a crucial role in improving model performance. Collecting more data could help the model better generalize, especially for the minority class (charged off).
 - Ensure that the data is representative of various customer profiles to make the model more robust.
6. Monitor Model Performance Over Time:
- Continuously monitor the model's performance as more data becomes available. Regular model updates might be necessary to adapt to changes in customer behavior or economic conditions.
7. Threshold Adjustment for Business Needs:
- Depending on business requirements, consider fine-tuning the threshold for classification based on the precision-recall curve. For example, if identifying defaulters is a priority, adjust the threshold to maximize recall at the expense of precision.
8. Deploy Model with Confidence:
- Given the solid performance of the Logistic Regression model (accuracy of 0.891), the model can be deployed with confidence for predictions on unseen data.
 - However, ensure that periodic reviews and updates are made to maintain accuracy and robustness.